

Software Verification

Master Programme in Computer Science

Mário Pereira mjp.pereira@fct.unl.pt

Nova School of Science and Technology, Portugal

November 17, 2025

Lecture 11

based on previous editions by João Seco, Luís Caires, and Bernardo Toninho
also based on lectures by Jean-Marie Madiot and Arthur Charguéraud

Verification of Heap-Manipulating Programs (3/n)

2. Concurrent ADTs
3. Concurrency Control via Monitors
4. Verifying Monitors
5. Conditions
6. Construction of Concurrent ADTs

Concurrent ADTs

To execute **several units** of computation out of the prescribed order (non strictly sequential).

An undetermined number of **interleaving executions** between units.

Allows **parallel computation** (Hardware permitting) – no state sharing.

Allows asynchronous programming (blocking IO)

- improves efficiency
- improves responsiveness (UI)

Examples of models with concurrency:

- Message passing
- Shared Memory
- Optimistic (Transactions)

To execute **several units** of computation out of the prescribed order (non strictly sequential).

An undetermined number of **interleaving executions** between units.

Allows **parallel computation** (Hardware permitting) – no state sharing.

Allows asynchronous programming (blocking IO)

- improves efficiency
- improves responsiveness (UI)

Examples of models with concurrency:

- Message passing
- **Shared Memory** – common in imperative setting
- Optimistic (Transactions)

Concurrency with Shared Memory

Several threads of execution **share** the same state footprint.

Interference is expected:

- the local view of a thread may change without notice (another thread may act “under the hood” and cause state changes).

Interference is **the essence** of concurrency.

Key issue:

- how to keep state consistency in the presence of state sharing and control interference
- how to reason about the effects of concurrent code

Reasoning about concurrency is **very challenging**.

Modularity brought by ADT based design is crucial!

The verification of sequential uses of ADTs characterizes the observable states of objects implementing them.

All ADT operations (**public** methods) preserve the consistency of the ADT.

Consistency is expressed by the **representation** and **abstract invariants** (and the abstraction mapping).

The following Hoare/Sep Logic triple is valid for all method bodies (`mbody`) of ADT operations

```
{ RepInv && pre-cond } mbody { RepInv && post-cond }
```

This line of reasoning works well under the assumption of **sequentiality**. What if method **executions overlap** in time?

Verification of Concurrent ADTs

Challenge: how to program and reason about ADTs with interfering methods.

Consider a Stack ADT

```
push(v)  
pop()  
isEmpty()
```

- push() interferes with pop()?
- pop() interferes with isEmpty()?
- pop() interferes with pop()?

Consider a Dictionary ADT

```
assoc(key, data)  
find(key)
```

- assoc() interferes with find()?
- assoc() interferes with assoc()?
- find() interferes with find()?

Verification of Concurrent ADTs

Challenge: how to program and reason about ADTs with interfering methods.

The verification of concurrent uses of ADTs should also observe stable states of objects implementing them.

All ADT operations (public methods) must also preserve the consistency of the ADT.

The following Hoare/Sep Logic triple should also be valid for all method bodies (mbody) of ADT operations

```
{ RepInv && pre-cond } mbody { RepInv && post-cond}
```

A sound approach is to consider the ADT operations as the unit of concurrency and structuring of reasoning (later to be refined).

Operation Level Behavior

Reason at the level of ADT operations and not at the level of unstructured low level instructions and state changes.

Each ADT operation is performed in three steps

- The operation is called (by the client thread)
- The operation is executed (inside the ADT)
- The operation returns

Example:

- push_call(2)
 - ... execute (internally to the ADT)
- push_return
- pop_call()
 - execute (internally to the ADT)
- pop_return(2)

Operation Level Behavior

We may consider several levels of concurrency.

- Several threads are invoking ADT operations but only one may actually be executing the operation.
 - strict serialisation, easier to implement and reason about
 - less chances of “unsound” interference

Several threads are invoking ADT operations but more than one may be executing an operation.

- more parallelism, more concurrency, harder to implement and reason about
- more chances of “unsound” / bad interference

How does the concurrent object behaviour relate to the intended sequential object specification ?

Serializability

- The global trace is always consistent with some sequential serialisation of previous operations (**no overlaps** of calls and returns), compatible with the sequential specification.

Linearizability

- The global trace is always consistent with a view in which previous operations appear to occur instantaneously between calls and returns, and the obtained serialisation is compatible with the sequential specification.

Linearizability is more flexible than serializability, as it allows for more parallel behaviour.

Desired properties ADT operations (ACID)

Atomicity

- No intermediate states are visible
(not compatible with the representation invariant)

Consistency

- Operations lead from a sound state to a sound state
(invariants and soundness are preserved)

Isolation

- This is another word for “no unsafe interference”

Durability

- Effects are undoable (important in the context of database transactions)

Correctness of Concurrent ADTs

With naive concurrency, it is hard (or impossible) for client code to be sure if a specific **post-condition** holds.

E.g.: two clients modify the concrete state at the same time, bringing the state inconsistent, breaking the representation invariant, or even crashing the code.

```
void push(int v)
//@ requires StackInv(this,?n,?m) && n < m;
//@ ensures StackInv(this,n+1,m);
{
    int i = nelems;
    store[i] = v;
    nelems++;
}

// Thread 1
int i = nelems;
store[i] = v;
nelems = i + 1;

// Thread 2
int i = nelems;
store[i] = v;
nelems = i + 1;
```

Correctness of Concurrent ADTs

With naive concurrency, it is hard (or impossible) for client code to be sure if a specific **post-condition** holds.

E.g.: two clients modify the concrete state at the same time, bringing the state inconsistent, breaking the representation invariant, or even crashing the code.

Solution using serialisation:

- **serialise** usages of concrete states, so that just a **single** thread may be accessing the state at each given moment (mutual exclusion of concrete state)
- We may then safely reason about such mutually exclusive code fragments as we have done for sequential code.

Correctness of Concurrent ADTs

With naive concurrency, it is hard (or impossible) for client code to be sure if a specific **post-condition** holds.

E.g.: two clients modify the concrete state at the same time, bringing the state inconsistent, breaking the representation invariant, or even crashing the code.

```
void pop()          // Thread 1          // Thread 2
/*@ requires StackInv(this,?n,?m)
   &*& n > 0; */
//@ ensures StackInv(this,n-1,m);
{
    int v = store[nelems-1];      int v = store[nelems-1];
    nelems--;                     nelems--;
    return v;                   int v = store[elements-1];
}                           nelems--;
                           return v;
                           }
                           return v;
```

Is this safe in all situations?

Correctness of Concurrent ADTs

With naive concurrency, it is hard (or impossible) for client code to be sure if a specific **pre-condition** holds.

E.g.: client checks that a buffer is not empty, but other thread empties it under the hood.

Solution:

- Concurrency control **replaces** pre-condition checking (on the client side) by **explicit waiting** for the precondition to hold (inside the ADT).
- The pre-condition for some ADT operation can only be enabled by **executing some other ADT operation**.
- So **waiting** for a pre-condition must be **managed** by special **programming language** or **system support**, in a coordinated way with other ADT operations.

Reasoning about concurrency is hard.

Making sure the code is right is much more difficult than in sequential code.

Trying to simulate the program running in your head and debugging is not really an option here.

We will now study how to design and construct correct concurrent code, based on **monitors**.

Monitor = invariant preserving concurrent ADT.

Nicely supported by `java.concurrent.util`

Concurrency Control via Monitors

An ADT where **operations** may be called **concurrently**.

2 key mechanisms provided for ensuring consistency:

- **Synchronization** (a.k.a. mutual exclusion)
 - only a single thread may “own” the shared state at any time object, and has permission to change it
 - all that client code may expect from shared state is the invariant, and **nothing more than the invariant**
 - any context switches must preserve the invariant (observable states)
- **Concurrency control**
 - pre-condition checking must be usually replaced by explicit waiting for the pre-condition to hold.
 - conditions refine the invariant into finer partitions.

Implementation of Monitors

To implement monitors in Java, we will use **locks**.

- A lot harder to reason about programs if we just think of using locks in an unstructured way.

We may later refine the borders of serialisability to get more concurrency (approach **linearisability**)

- Use locks as delimiters of abstract operations on the shared state
- Use the `java.util.concurrent` API (Doug Lea)
- Will learn how to design concurrent ADTs without thinking “operationally”, but rather in terms of (partitioned) ownership, invariants, and conditions.

Example – Bounded Counter

Consider a counter with a maximum value:

```
/*@
 predicate BCounterInv(BCounter c; int v,int m) = c.N |-> v &*& c.MAX |-> m &*&
           v>=0 &*& v<=m;
*/
class BCounter {
    int N; int MAX;

    BCounter (int max)
    // @ requires 0 <= max;
    // @ ensures BCounterInv(this,0,max);
    { N = 0; MAX = max; }

    void inc ()
    // @ requires BCounterInv(this,?n,?m) &*& n < m;
    // @ ensures BCounterInv(this,n+1,m);
    { N++; }

    void dec ()
    // @ requires BCounterInv(this,?n,?m) &*& n > 0;
    // @ ensures BCounterInv(this,n-1,m);
    { N--; }

    int get ()
    // @ requires BCounterInv(this,?n,?m);
    // @ ensures BCounterInv(this,n,m) &*& 0<=result &*& result<=m;
    { return N; }
}
```

Example – Bounded Counter

The use of the bounded counter is safe in a sequential setting.

```
public static void main(String[] args)
//@ requires true;
//@ ensures true;
{
    int MAX = 100;
    BCounter c = new BCounter(MAX);
    //@ assert BCounterInv(c,0,MAX);
    if (c.get() < MAX) {
        c.inc(); // this is ok, precondition satisfied
    }
}
```

Example – Bounded Counter

What happens in a concurrent setting where the reference to the counter is used elsewhere?

```
public static void main(String[] args)
//@ requires true;
//@ ensures true;
{
    int MAX = 100;
    BCounter c = new BCounter(MAX);
    //@ assert BCounterInv(c,0,MAX);
    giveAway(c); // potentially give other thread access to c
    if (c.get() < MAX) {
        //@ assert BCounterInv(c,?v,MAX) &*& v < MAX;
        c.inc();
        // not safe any more as other thread may have acted
    }
}
```

Serialise access to shared state

```
import java.util.concurrent.*;
import java.util.concurrent.locks.*;

class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;

    CCounter(int max) {
        N = 0;
        MAX = max;
        mon = new ReentrantLock();
    }
    ...
}
```

Example – Bounded Counter

Consider a counter with a maximum value:

```
import java.util.concurrent.*;
import java.util.concurrent.locks.*;

class BCounter {
    int N; int MAX;
    ReentrantLock mon;

    void inc ()
    {
        mon."enter"(); //request permission to the shared state
        N++;
        mon."leave"(); //release ownership of the shared state
    }

    void dec ()
    {
        mon."enter"(); //request permission to the shared state
        N--;
        mon."leave"(); //release ownership of the shared state
    }
}
```

Example – Bounded Counter

Consider a counter with a maximum value:

```
import java.util.concurrent.*;
import java.util.concurrent.locks.*;

class BCounter {
    int N; int MAX;
    ReentrantLock mon;

    void inc ()
    {
        mon.lock(); //request permission to the shared state
        N++;
        mon.unlock(); //release ownership of the shared state
    }

    void dec ()
    {
        mon.lock(); //request permission to the shared state
        N--;
        mon.unlock(); //release ownership of the shared state
    }
}
```

Example – Bounded Counter

Consider a counter with a maximum value:

```
import java.util.concurrent.*;
import java.util.concurrent.locks.*;

class BCounter {
    int N; int MAX;
    ReentrantLock mon;

    int get ()
    {
        int r;
        mon."enter"();
        r = N; // put a copy on the stack, private to the thread
        mon."leave"();
        return r;
    }
}
```

Undisciplined concurrent access to shared memory causes unexpected (thus invalid) behaviour.

(Conservatively) Implement all operations in a monitor, where only one thread is allowed at a time.

```
T op()
{
    T result;
    mon. "enter"();
    // Thread has access to shared state
    result = Expression;
    // Thread releases access to shared state
    mon. "leave"();
    return result;
}
```

(Later) we will define more flexible mechanisms that allow more interleaving of operations.

Verifying Monitors

Reasoning About Monitor Operations

Monitor “guards” the access to the ADT shared state.

shared state = footprint of the monitor operations

The “**enter**” operation gives access to the footprint.

The “**leave**” operation captures the footprint back.

Hoare/Separation Rule for lock/unlock

{emp} $m.\text{enter}()$ {SharedStateInv()}

{SharedStateInv()} $m.\text{leave}()$ {emp}

SharedStateInv() is the representation invariant of the ADT that is used to verify operations in exclusive ownership of the monitor.

In our example

```
/*@ predicate CCounterInv(CCounter c) =
  c.N |-> ?v && c.MAX |-> ?m && v>=0 && v<=m; @*/
```

so:

{emp} $m.\text{lock}()$ {CCounterInv(this)}

Hoare/Separation Rule for lock/unlock

Representation Invariant is available inside the monitor.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;
    void inc()
    //@ requires ...
    //@ ensures ...
    {
        //@ request permission to the shared state
        mon.lock();
        //@ assert CCounterInv(this,?n,?m)
        N++;
        //@ assert CCounterInv(this,n+1,m)
        mon.unlock();
        //@ release ownership of the shared state
    }
}
```

Representation Invariant is available inside the monitor.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;
    void inc()
    //@ requires CCounterInv(this,?n,?m) &*& n < m;
    //@ ensures CCounterInv(this,n+1,m);
    {
        //@ request permission to the shared state
        mon.lock();
        //@ assert CCounterInv(this ,?n ,?m)
        N++;
        //@ assert CCounterInv(this ,n+1 ,m)
        mon.unlock();
        //@ release ownership of the shared state
    }
}
```

... assertions in red are not available in pre- or post-conditions.

But, What Happens Then?

Many questions arise now:

- How can a client check $n < m$?
- How can an ADT capture the footprint in a lock?

The “Invisible” Abstract State

Many threads may be interfering, so the only thing one may assume is the **invariant**, only after **entering** the **shared state** a client may know extra details about the concrete state.

In fact, nothing **specific** about the **abstract state** may be revealed to client code, and we need to be **less informative** about the abstract state (e.g., no current val).

Inside the object, the **only unprotected** objects are the **locks** (or the single lock).

Each lock can be used to ask **permission** to access a disjoint part of the shared state.

We must precisely define **which part** of the shared state is **separately owned** by each lock.

“Invisible” Abstract State

```
/*@
 predicate_ctor CCounter_shared_state (BCounter c) () =
    c.N |-> ?v && v >= 0 && c.MAX |-> ?m && 0 < m
    && v <= m;
*/
/*@ predicate CCounterInv(CCounter c) =
    c.mon |-> ?l
    && l != null
    && lck(l, 1, CCounter_shared_state(c))
*/
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;

    CCounter(int max)
    //@ ensures CCounterInv(this);
    {
        N = 0 ;
        MAX = max;
        mon = new ReentrantLock();
    }
}
```

Higher-order Predicates

Define a family of predicates that can be stored by other predicates through a **predicate constructor**

```
/*@
  predicate_ctor CCounter_shared_state (BCounter c) () =
    c.N |-> ?v &*& v >= 0 &*& c.MAX |-> ?m &*& 0 < m
    &*& v <= m;
@*/
```

and then instantiate them to produce a predicate instance

```
//@ lck(1, 1, CCounter_shared_state(c))
```

and use it in other predicates

```
//@ CCounter_shared_state(this)()
```

ReentrantLock Constructor

The ReentrantLock class signature ensures the establishment of the “native” predicate `lck` based on the invariant predicate `inv`

```
/*@ predicate enter_lck(real p, predicate() inv) = inv();
   predicate lck(Lock s; real p, predicate() inv); @*/
public class ReentrantLock {
    public ReentrantLock();
    //@ requires enter_lck(1,?inv);
    //@ ensures lck(this, 1, inv);
    ...
}
```

Example – Bounded Counter

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;

    public CCounter (int max)
    //@ ensures CCounterInv(this);
    {
        N = 0 ;
        MAX = max;
        //@ close CCounter_shared_state(this)();
        //@ close enter_lck(1, CCounter_shared_state(this));
        mon = new ReentrantLock();
        //@ assert lck(mon, 1, CCounter_shared_state(this));
        //@ close CCounterInv(this);
    }
}
```

Example – Bounded Counter

Representation invariant captures the resources that can be used in a concurrent context.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;

    public void inc()
    //@ requires CCounterInv(this);
    //@ ensures CCounterInv(this);
    {
        //@ open CCounterInv(this);
        mon.lock(); // request permission to the shared state
        //@ open CCounter_shared_state(this)();
        N++;
        //@ close CCounter_shared_state(this)();
        mon.unlock(); // release ownership of the shared state
    }
}
```

ReentrantLock Operations

The ReentrantLock methods `lock` and `unlock` use the predicate `lck` to make the invariant available inside the monitor and captured again outside the exclusive access region.

```
/*@
 predicate lck(Lock s; real p, predicate() inv);
@*/

public class ReentrantLock {
    ...
    public void lock();
    //@ requires lck(?t, 1, ?inv);
    //@ ensures lck(t, 0, inv) && inv();

    public void unlock();
    //@ requires lck(?t, 0, ?inv) && inv();
    //@ ensures lck(t, 1, inv);
    ...
}
```

Example – Bounded Counter

Representation invariant captures the resources that can be used in a concurrent context.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;

    public void dec()
    //@ requires CCounterInv(this);
    //@ ensures CCounterInv(this);
    {
        //@ open CCounterInv(this);
        mon.lock(); // request permission to the shared state
        //@ open CCounter_shared_state(this)();
        N--;
        //@ close CCounter_shared_state(this)();
        mon.unlock(); // release ownership of the shared state
    }
}
```

The “public” representation invariant **cannot reveal** information to make the **pre-condition hold**.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;

    void dec()
    // this is the only pre- one can write
    //@ requires CCounterInv(this);
    //@ ensures CCounterInv(this);
    {
        //@ open CCounterInv(this);
        mon.enter();
        //@ open CCounter_shared_state(this)();
        N--;
        // must ensure N >= 0
        //@ close CCounter_shared_state(this)();
        mon.leave();
    }
}
```

This structured monitor implementation makes the shared state representation invariant available inside the exclusive access area.

The operations `lock` and `unlock` release and capture the **representation invariant**.

Operations **cannot enforce preconditions** in the caller context due to possible interleaving of operations outside the monitor.

Conditions

Lets get back to the “difficult” questions:

- How can a client check $n < m$?
- What if $N == 0$?

How can a client check $n < m$?

Concurrency control mechanisms replace pre-condition checking (on the client side) by explicit waiting for the precondition to hold (inside the ADT).

The pre-condition for some ADT operation can only be enabled by executing some other ADT operation.

Waiting for a pre-condition is managed by special programming language or system support, in a coordinated way with other ADT operations.

Monitor Conditions

Partition Shared State Using Conditions

Conditions implement queues of suspended threads.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;
    Condition notZero;
    Condition notMax;

    public void dec()
    //@ requires CCounterInv(this);
    //@ ensures CCounterInv(this);
    {
        //@ open CCounterInv(this);
        mon.lock(); // request permission to the shared state
        //@ open CCounter_shared_state(this)();
        if (N == 0) notZero.await();
        N--;
        //@ close CCounter_shared_state(this)();
        mon.unlock(); // release ownership of the shared state
    }
}
```

Partition Shared State Using Conditions

Conditions represent operations' preconditions, that are checked in the callee context, where access is granted.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;
    Condition notZero;
    Condition notMax;

    public void dec()
    //@ requires CCounterInv(this);
    //@ ensures CCounterInv(this);
    {
        //@ open CCounterInv(this);
        mon.lock();
        //@ open CCounter_shared_state(this)();
        if (N == 0) notZero.await(); // check for condition to
            hold
        N--;
        //@ close CCounter_shared_state(this)();
        mon.unlock();
    }
}
```

Hoare/Separation Rule for await

$$\{ \text{SharedStateInv}() \} C.\text{await}() \{ \text{SharedStateInv} \star \text{Cond}(C) \}$$

$\text{Cond}(C)$ is the refinement of the shared state property denoted by C .

In the Bounded Counter example:

- $\text{Cond(notZero)} = N > 0$
- $\text{Cond(notMax)} = N < \text{MAX}$

Partition Shared State Using Conditions

```
/*@
predicate_ctor CCounter_shared_state (CCounter c) () =
  c.N |-> ?v && c.MAX |-> ?m && v >= 0 && m > 0 && v <= m;

predicate_ctor CCounter_notzero (CCounter c) () =
  c.N |-> ?v && c.MAX |-> ?m && v > 0 && m > 0 && v <= m;

predicate_ctor CCounter_notmax (CCounter c) () =
  c.N |-> ?v && c.MAX |-> ?m && v > 0 && m >= 0 && v < m;

predicate CCounterInv(CCounter c) =
  c.mon |-> ?l
  && l != null
  && lck(l,1, CCounter_shared_state(c))
  && c.notzero |-> ?cc
  && cc != null
  && cond(cc, CCounter_shared_state(c), CCounter_notzero(c))
  && c.notmax |-> ?cm
  && cm != null
  && cond(cm, CCounter_shared_state(c), CCounter_notmax(c));
@*/
```

ReentrantLock Conditions `await`

The Condition method `await` makes a transformation from the generic shared state (`inv`) to the refined state corresponding to that condition (`acond`).

```
public interface Condition {  
  
    public void await();  
    //@ requires cond(this,?inv,?acond) &*& inv();  
    //@ ensures cond(this, inv, acond) &*& acond();  
}
```

ReentrantLock Conditions await

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;
    Condition notZero;
    Condition notMax;

    public void dec()
    //@ requires CCounterInv(this);
    //@ ensures CCounterInv(this);
    {
        //@ open CCounterInv(this);
        mon.lock();
        //@ open CCounter_shared_state(this)();
        if (N == 0) notZero.await();
        // refined state with N > 0
        //@ open CCounter_notzero(this)();
        N--;
        //@ close CCounter_shared_state(this)();
        mon.unlock();
    }
}
```

Example – Bounded Counter

Conditions are defined with relation to both the representation invariant and the refined state.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;
    Condition notZero;
    Condition notMax;

    BCounter(int max)
    //@ requires max > 0;
    //@ ensures CCounterInv(this);
    {
        MAX = max;
        mon = new ReentrantLock();

        //@ close CCounter_shared_state(this)();
        //@ close set_cond(CCounter_shared_state(this),CCounter_notzero(this));
        notZero = mon.newCondition(); // notZero set to mean N > 0

        //@ close set_cond(CCounter_shared_state(this),CCounter_notmax(this));
        notMax = mon.newCondition(); // notMax set to mean N < MAX
    }
}
```

ReentrantLock Conditions

Conditions are defined with relation to **both** the

- representation invariant = **inv**
- and the refined state. = **pred**

```
public class ReentrantLock {  
    ...  
  
    public Condition newCondition();  
    // @ requires lck(?t, 1, ?inv) &*& set_cond(inv, ?pred);  
    /*@ ensures lck(t, 1, inv) &*& result != null  
        &*& cond(result, inv, pred); @*/  
}
```

Partition Shared State Using Conditions

Different conditions give access to different refinements.

```
class CCounter {
    int N;
    int MAX;
    ReentrantLock mon;
    Condition notZero; Condition notMax;

    void inc()
    //@ requires CCounterInv(this);
    //@ ensures CCounterInv(this);
    {
        //@ open CCounterInv(this);
        mon.lock();
        //@ open CCounter_shared_state(this)();
        if (N == MAX) notmax.await();
        //@ open CCounter_notmax(this)(); // refined state N < MAX
        N++;
        //@ close CCounter_shared_state(this)();
        mon.unlock();
    }
}
```

Wait and Signal Work Together

Wait and signal represents a voluntary yielding of control flow and exclusive access to the monitor.

Wait **suspends** a thread until a precondition is **satisfied**.

Signal **states** that a precondition is explicitly **satisfied**.

Ensure Progress Using Signaling

```
void inc()
//@ requires CCounterInv(this);
//@ ensures CCounterInv(this);
{
    //@ open CCounterInv(this);
    mon.lock();
    //@ open CCounter_shared_state(this)();
    if (N == MAX) notMax.await();
    //@ assert CCounter_notmax(this)();
    N++;
    //@ close CCounter_notzero(this)();
    notZero.signal();
    //@ close CCounter_shared_state(this)();
    mon.unlock();
}
```

Ensure Progress Using Signaling

```
void dec()
//@ requires CCounterInv(this);
//@ ensures CCounterInv(this);
{
    //@ open CCounterInv(this);
    mon.lock();
    //@ open CCounter_shared_state(this)();
    if (N == 0) notZero.await();
    //@ assert CCounter_notzero(this)();
    N--;
    //@ close CCounter_notmax(this)();
    notMax.signal();
    //@ close CCounter_shared_state(this)();
    mon.unlock();
}
}
```

Hoare/Separation Rule for await/signal

$\{ \text{SharedStateInv}() \} C.\text{await}() \{ \text{SharedStateInv}() \star \text{Cond}(C) \}$

$\{ \text{SharedStateInv}() \star \text{Cond}(C) \} C.\text{signal}() \{ \text{SharedStateInv}() \}$

$\text{Cond}(C)$ is the refinement of the shared state property denoted by condition C.

In the Bounded Counter example:

- $\text{Cond(notZero)} = N > 0$
- $\text{Cond(notMax)} = N < \text{MAX}$

Defending Against Unsound Implementation

Excerpt from Java API documentation:

Implementation Considerations

When waiting upon a Condition, a "spurious wakeup" is permitted to occur, in general, as a concession to the underlying platform semantics. This has little practical impact on most application programs as a Condition **should always be waited upon in a loop**, testing the state predicate that is being waited for.

An implementation is free to remove the possibility of spurious wakeups but it is recommended that applications programmers always assume that they can occur and so always wait in a loop.

Defending Against Unsound Implementation

```
void inc()
//@ requires CCounterInv(this);
//@ ensures CCounterInv(this);
{
    //@ open CCounterInv(this);
    mon.lock();
    //@ open CCounter_shared_state(this)();
    while (N == MAX) notMax.await();
    //@ assert CCounter_notmax(this)();
    N++;
    //@ close CCounter_notzero(this)();
    notZero.signal();
    //@ close CCounter_shared_state(this)();
    mon.unlock();
}
```

Defending Against Unsound Implementation

```
void dec()
//@ requires CCounterInv(this);
//@ ensures CCounterInv(this);
{
    //@ open CCounterInv(this);
    mon.lock();
    //@ open CCounter_shared_state(this)();
    while (N == 0) notZero.await();
    //@ assert CCounter_notzero(this)();
    N--;
    //@ close CCounter_notmax(this)();
    notMax.signal();
    //@ close CCounter_shared_state(this)();
    mon.unlock();
}
```

Defending Against Unsound Implementation

```
void inc()
//@ requires CCounterInv(this);
//@ ensures CCounterInv(this);
{
    //@ open CCounterInv(this);
    mon.lock();
    //@ open CCounter_shared_state(this)();
    while (N == MAX)
        /*@ invariant this.N |-> ?v &*& v >= 0
           &*& this.MAX |-> ?m
           &*& m > 0 &*& v <= m
           &*& this.notzero |-> ?cc &*& cc !=null
           &*& cond(cc,CCounter_shared_state(this), CCounter_nonzero(this))
           &*& this.notmax |-> ?cm
           &*& cm !=null
           &*& cond(cm, CCounter_shared_state(this),CCounter_nonmax(this));
        */
    {
        //@ close CCounter_shared_state(this)();
        try { notMax.await(); } catch (InterruptedException e) {}
        //@ open CCounter_notmax(this)();
    }
    N++;
    //@ close CCounter_notzero(this)();
    notZero.signal();
    //@ close CCounter_shared_state(this)();
    mon.unlock();
    //@ close CCounterInv(this);
}
```

Implementing preconditions using monitor conditions.

The operations `await` and `signal` yield control in a structured way.

Protect against unsound implementations: wait inside a loop.

Construction of Concurrent ADTs

Concurrent ADT Construction Steps

Challenge: Can we systematically transform and verify correct a “sequential” ADT implementation into an efficient “concurrent” ADT implementation?

Concurrent ADT Recipe

1. Associate a monitor to the ADT (ReentrantLock - `mon`)
2. Define the **sequential ADT Representation Invariant** (`RepInv`) that talks about the shared state
3. Define the **concurrent ADT Representation Invariant** that talks about the **monitor** and **associated conditions**.
4. In the implementation of each operation of the CADT
 - 4.1 get access to the shared state representation invariant, use `mon.lock()`
 - 4.2 when done and only if the shared state representation invariant holds, use `mon.unlock()`
5. Replace the **ADT operation pre-conditions** by **monitor conditions** inside the monitor.
6. Design voluntary **yielding** points to implement the correct interleaving of operations.

Concurrent ADT Construction Steps

To replace ADT operation pre-conditions by monitor conditions inside the monitor, we must consider the following aspects:

- When a thread enters a CADT operation and gets ownership of the representation invariant, it should check the state to satisfy (or not) the pre-condition (e.g., wants to decrement but counter value is zero).
- The thread must then await for the condition to hold (e.g., for the value to be > 0).
- Conversely, whenever a thread running inside the CADT establishes any one of the monitor conditions (e.g., inc establishes value > 0), it has the duty to signal the condition (so that the runtime system may awake a waiting thread).
- Note that signaling is there to **help the system to progress**, and simplify the implementation of monitors.